

A deep dive on ERC-20 contract vulnerabilities

No Author Given

No Institute Given

Abstract. ERC-20 is the most prominent Ethereum standard for transferable tokens. Tokens implementing the ERC-20 interface can interoperate with a large number of already deployed internet-based services and Ethereum-based smart contracts. In recent years, security vulnerabilities in ERC-20 implementations have been uncovered. We systemize these across 7 auditing tools into a set of 82 distinct vulnerabilities and best practices. Next, we use our experience to provide a new secure implementation of the ERC-20 interface, **TokenHook**, that is freely¹ available. Reference ERC-20 implementations have been slowly abandoned over time. **TokenHook** has enhanced security properties and stronger compliance with best practices compared to the sole surviving reference implementation (from OpenZeppelin) in the ERC-20 specification.

1 Introduction

Ethereum blockchain [18,74] allows users to build and deploy decentralized applications (DApps) that can accept and use ETH as its protocol-level cryptocurrency. Many DApps also issue their own custom tokens with a variety of intents, including tokens as: financial products, in-house currencies, voting rights for DApp governance, valuable assets, crypto-collectibles, *etc.* To encourage interoperability with other DApps and web apps (exchanges, wallets, *etc.*), the Ethereum community accepted a popular token standard (for non-fungible tokens) called ERC-20 [23]. While numerous ERC-20 extensions or replacements have been proposed, ERC-20 remains prominent. Of the 2.5M[53] smart contracts on the Ethereum network, 260K are tokens [66]. 98% of these tokens are ERC-20 [21], demonstrating their widespread acceptance by the industry and Ethereum community.

The development of smart contracts has been proven to be error-prone, and as a result, smart contracts are often riddled with security vulnerabilities. An early study in 2016 found that 45% of smart contracts at that time had vulnerabilities[38]. ERC-20 token are subset of smart contracts and security is particularly important given that many tokens have considerable market capitalization (*e.g.*, USDT, LINK, CRO, LEO, DAI, *etc.*). As tokens can be held by commercial firms, in addition to individuals, and firms need audited financial statements in certain circumstances, the correctness of the contract issuing the tokens is now in the purview of professional auditors. One tool we examine, EY Smart Contract and Token Review[22], is from a ‘big-four’ auditing firm.

¹ Implementation on Etherscan with source code and deployed on Mainnet and Rinkeby: <https://bit.ly/35FMbAf>, <https://bit.ly/33wDENx>

Contributions. Similar to any new technology, Ethereum has undergone numerous security attacks that have collectively caused more than US\$100M in financial losses[25,46,44,58,49,3]. Although research has been done on smart contract vulnerabilities in the past[32], our focus is on ERC-20 tokens only. Some vulnerabilities (such as multiple withdrawals) will be more serious in token contracts. This motivates us to (i) comprehensively study all known vulnerabilities in ERC-20 token contracts, systematizing them into a set of 82 distinct vulnerabilities and best practices, and review the completeness and precision of auditing tools in detecting these vulnerabilities to establish the reliability of an audit based on these tools. We (ii) use this research to provide a new secure implementation of the ERC-20 interface, **TokenHook**, that is freely available and open source. Compared to other implementations from OpenZeppelin[42] and ConsenSys[8], it is more secure and fully compatible with ERC-20 specifications. Finally, (iii) we examine the practicality of our work in the context of the top ten ERC-20 tokens by market capitalization.

2 TokenHook

TokenHook is our ERC20-compliant implementation written in Solidity. TokenHook is open source and available on Etherscan, where it has been tested with MetaMask and deployed on Mainnet.² It can be customized by developers, who can refer to each mitigation technique separately and address a specific attack. Required comments have been also added to clarify usage of each part. Standard functionalities of the token (*i.e.*, `approve()`, `transfer()`, *etc.*) have been unit tested. A demonstration of token interactions and event triggering can also be seen on Etherscan.³

Among the layers of the Ethereum blockchain, ERC-20 tokens fall under the *Contract layer* in which DApps are executed. The presence of security vulnerability in supplementary layers affect the entire Ethereum blockchain, not necessarily ERC-20 tokens. Therefore, vulnerabilities in other layers are assumed to be out of the scope (*e.g.*, *Indistinguishable chains* at the data layer, the *51% attack* at the consensus layer, *Unlimited nodes creation* at network layer, and *Web3.js Arbitrary File Write* at application layer). Moreover, we exclude vulnerabilities identified in now outdated compiler versions, for example:

- *Constructor name ambiguity* in versions before 0.4.22.
- *Uninitialized storage pointer* in versions before 0.5.0.
- *Function default visibility* in versions before 0.5.0
- *Typographical error* in versions before 0.5.8.
- *Deprecated solidity functions* in versions before 0.4.25.
- *Assert Violation* in versions before 0.4.10.
- *Under-priced DoS attack* before EIP-150 & EIP-1884.

² Etherscan: <https://bit.ly/35FMbAf>

³ Etherscan: <https://bit.ly/33xHfL2>, <https://bit.ly/35TimMW>

2.1 Security features

In Appendix A, we examine general attack vectors and cross-check their applicability to ERC-20 tokens. As many of these are now well-researched attacks, we leave them in the appendix. How **TokenHook** mitigates these attacks is summarized as follows.

1. We secure the `transferFrom()` function to mitigate the *multiple withdrawal* attack [51]. Without our counter-measure, an attacker can use a front-running attack [9,17] to transfer more tokens than what is intended (approved) by the token holder. Our solution is compliant with the ERC-20 standard. (*cf.* Appendix A.1)
2. We use the `SafeMath` library in all arithmetic operations to catch over/under flows. (*cf.* Appendix A.2)
3. We implement a `noReentrancy` modifier for external functions to mitigate *same-function re-entrancy* and *cross-function re-entrancy* attacks using mutual exclusion (mutex). (*cf.* Appendix A.3)
4. We check the return value of `call.value()` to revert failed fund transfers in `sell()` and `withdraw()` functions. It mitigates the *unchecked return values* attack while making the token contract compatible with EIP-1884 [31]. (*cf.* Appendix A.4)
5. We mitigate the *frozen ether* issue by defining a `withdraw()` function that allows the owner to transfer all ETH out of the token contract. Otherwise, unexpected ETH forced onto the token contract (*e.g.*, from another contract running `selfdestruct`) will be stuck forever (*cf.* Appendix A.5)
6. We apply an `onlyOwner` modifier to the `withdraw()` function to the mitigate *unprotected Ether withdrawal* issue by enforcing authentication before transferring any funds out of the contract. (*cf.* Appendix A.6)
7. We use embedded `Library` code to reduce gas costs (calling functions in embedded libraries requires less gas than external calls) and mitigate the *state variable manipulation* attack. (*cf.* Appendix A.7)
8. We carefully define the visibility of each function. Most of the functions are declared as `External` (*e.g.*, `Approve()`, `Transfer()`, *etc.*) per specifications of ERC-20 standard. (*cf.* Appendix A.8)

2.2 Best practices and enhancements

In Appendix B, we also take into account a number of best practices for developing DApps and discuss those that are applicable to ERC-20. What follows is an overview of how we implement these in **TokenHook**.

1. We implement all ERC-20 functions to make it fully compatible with the standard. Compliance is important for ensuring that other DApps and web apps (*i.e.*, crypto-wallets, crypto-exchanges, web services, *etc.*) compose with **TokenHook** as expected. (*cf.* Appendix B.1)
2. We apply an `external` visibility for interactive functions (*e.g.*, `approve()` and `transfer()`, *etc.*) to improve performance. External functions can read

arguments directly from non-persistent `calldata` instead of allocating persistent memory by the EVM. (*cf.* Appendix B.2)

3. We implement a ‘cease trade’ operation that will freeze the token in the case of new security threats or new legal requirements (*e.g.*, Liberty Reserve [73] or TON cryptocurrency[16]). To freeze all functionality of `TokenHook`, the owner (or multiple parties) can call the function `pause()` which sets a lock variable. All critical methods are marked with a `notPaused` modifier that will throw exceptions until functionality is restored using `unpause()`. (*cf.* Appendix B.3)
4. We define nine extra events: `Buy`, `Sell`, `Received`, `Withdrawal`, `Pause`, `Change`, `ChangeOwner`, `Mint` and `Burn`. `Change` event logs any state variable updates that can be used to watch for token inconsistent behaviour (*e.g.*, via `TokenScope`[5]) and react accordingly. (*cf.* Appendix B.4)
5. We implement a `sell()` and `buy()` function for exchanging between tokens and ETH. `sell()` allows token holders to exchange tokens for ETH and `buy()` accepts ETH by adjusting buyer’s token balance. While this is not necessary under ERC-20, we have seen this functionality added to tokens (*e.g.*, ORBT [52]) and it needs to be done securely to prevent attacks like re-entrancy. It is used to buy and sell tokens at a fixed price (*e.g.*, an initial coin offering (ICO), prediction market portfolios [6]) independent of crypto-exchanges, which introduce a delay (for the token to be listed) and fees.
6. We choose to make `TokenHook` non-upgradable so it can be audited, and upgrades will not introduce new vulnerabilities that did not exist at the time of the initial audit.
7. We also follow other best practices such as not using batch processing in `sell()` function to avoid *DoS with unexpected revert* issue, not using miner controlled variable in conditional statements, and not using `SELFDESTRUCT`.

2.3 Need for another reference implementation

The authors of the ERC-20 standard reference two sample Solidity implementations: one that is actively maintained by OpenZeppelin [42] and one that has been deprecated by ConsenSys [8] (and now refers to the OpenZeppelin implementation). As expected, the OpenZeppelin template is very popular within the Ethereum community [55,75,48].

Diversity in software is important for robustness and security [26,27]. For ERC-20, a variety of implementations will reduce the impact of a single bug in a single implementation. For example, between 17 March 2017 and 13 July 2017, OpenZeppelin’s implementation used the wrong interface and affected 130 tokens [10]. This is our primary motivation for developing `TokenHook`.

OpenZeppelin’s implementation is actually part of a small portfolio of implementations (ERC20, ECR721, ERC777, and ERC1155). Code reuse across the four implementations adds complexity for a developer that only wants ERC20. Further, most audit tools are not able to import libraries/interfaces from external files (*e.g.*, `SafeMath.sol`, `IERC20.sol`). By contrast, `TokenHook` uses a flat layout in a single file that is specific to ERC20.

TokenHook makes other improvements over the OpenZeppelin implementation. OpenZeppelin introduces two new functions to mitigate the multiple withdraw attack: `increaseAllowance()` and `decreaseAllowance()` however these are not part of the ERC-20 standard and are not interoperable with other applications that expect to use `approve()` and `transferFrom()`. TokenHook secures `transferFrom()` to prevent the attack (following [51]) and is interoperable with legacy DApps and web apps. Additionally, TokenHook mitigates the *frozen Ether* issue by introducing a `withdraw()` function, while ETH forced into the OpenZeppelin implementation is forever unrecoverable. Both contracts implement a *fail-safe mode*, however this logic is internal to TokenHook, while OpenZeppelin requires an external `Pausable.sol` contract.

In terms of best practices, OpenZeppelin requires other optimizations such as, locking the *pragma* to avoid using outdated compiler versions by developers (*cf.* SWC-103), emit a *Change* event (*cf.* TokenScope[5]) when updating state variables (*e.g.*, `_decimals=decimals_` in `setupDecimals()`). Initializing *totalSupply* in the constructor is also recommended. As we implemented, using `external` visibility instead of `public` increases readability and consumes less gas. Finally, these are very minor but it helps to have more readable token contract by avoiding similar variable names (*e.g.*, `_name=name_` in `constructor()`) and using *mixedCase* format when declaring variable and functions.

3 Code audit

We used a variety of code audit tools on TokenHook to validate the code and also to illuminate the completeness and error-rate of such tools on one specific use-case (similar work studies in less depth a variety of use-cases[2]). We also did not adapt older tools that support significantly lower versions of the Solidity compiler (*e.g.*, Oyente). We provide a version number for software tools; the remaining tools are web-based and were used in the summer of 2020:

1. EY Review Tool by Ernst & Young Global Limited[22].
2. SmartCheck by SmartDec[62].
3. Securify v2.0 by ChainSecurity[68].
4. ContractGuard by GuardStrike[29].
5. MythX by ConsenSys[7].
6. Slither Analyzer v0.6.12 by Crytic[36].
7. Odin by Sooho[63].

3.1 Analysis of audit results

A total of 82 audits have been conducted by these auditing tools. Audits include best practices and security vulnerabilities. The results are summarized in Tables1–3. To compile the list, we referenced the knowledge-base of each tool[67,62,7,29,36], understood each threat, manually mapped the audit to the corresponding SWC registry[61], and manually determined when different tools were testing for the same vulnerability or best practice (which was not always

Table 1. Auditing results of 7 smart contract analysis tools on TokenHook. ✓=Passed audit, ⊕=False positive, ×=Failed audit, Empty=Not supported audit by the tool, !=Informational, ○=Tool specific audit (No SWC registry), BP=Best practice

ID	SWC	Vulnerability or best practice Mitigation or recommendation	Security tools					
28	127	Arbitrary Jump with Function Type Variable Minimizing use of assembly in the code	✓	✓	✓	✓	✓	✓
29	128	DoS With Block Gas Limit Avoiding loops across the code that may consume considerable resources	✓	✓	✓	✓	✓	✓
30	129	Typographical Error Using SafeMath library or performing checks on any math operation			✓			✓
31	130	Right-To-Left-Override control character (U+202E) Avoiding U+202E character which forces RTL text rendering		✓	✓	✓	✓	✓
32	131	Presence of unused variables Removing all unused variables to decrease gas consumption	✓	✓		✓	✓	⊕
33	132	Unexpected Ether balance Avoiding Ether balance check in the code (<i>e.g.</i> , this.balance == 0.24 Ether)	✓	✓		✓	✓	✓
34	133	Hash Collisions With Variable Length Arguments Using abi.encode() instead of abi.encodePacked() to prevent hash collision						✓
35	134	Message call with hardcoded gas amount Using .call.value("") which is compatible with EIP1884	⊕	⊕	✓	✓		✓
36	135	Code With No Effects Writing unit tests to ensure producing the intended effects by DApps	✓					✓
37	136	Unencrypted Private Data On-Chain Storing un-encrypted private data off-chain	!					✓
38	○	Allowance decreases upon transfer Decreasing allowance in transferFrom() method	✓					
39	○	Allowance function returns an accurate value Returning only value from the mapping instead of internal function logic	✓					
40	○	It is possible to cancel an existing allowance Possibility of setting allowance to 0 to revoke previous allowances	✓	✓				
41	○	A transfer with an insufficient amount is reverted Checking balances in transfer() method before updating balances	✓				✓	
42	○	Upon sending funds, the sender's balance is updated Updating balances in transfer() or transferFrom() methods	✓					
43	○	The Transfer event correctly logged Emitting Transfer event in transfer() or transferFrom() functions	✓					
44	○	Transfer an amount that is greater than the allowance Checking balances in transferFrom() method before updating balances	✓					
45	○	Risk of short address attack is minimized Using recent Solidity version to mitigate the attack	✓			✓		
46	○	Function names are unique No function overloading to avoid unexpected behavior	✓				✓	
47	○	Using miner controlled variables Avoiding block.number, block.timestamp, block.difficulty, now, etc	✓	✓	✓	✓	✓	✓
48	○	Use of return in constructor Not using return in contract's constructor	✓					
49	○	Throwing exceptions in transfer() and transferFrom() Returning true after successful execution or raising exception in failures	✓				✓	
50	○	State variables that could be declared constant Adding constant attribute to variables like name, symbol, decimals, etc					✓	
51	○	Tautology or contradiction Fixing comparison in the code that are always true or false					✓	
52	○	Divide before multiply Ordering multiplication prior division to avoid integer truncation					✓	
53	○	Unchecked Send Ensuring that the return value of send() is always checked					✓	
54	BP	Too many digits Using scientific notation to make the code readable and simpler to debug					✓	

Table 2. Continuation of Table1.

			EY Token Review Smart Check Security MythX (Mythril) Contract Guard Slither Odin						
ID	SWC	Vulnerability or best practice Mitigation or recommendation	Security tools						
55	BP	The decreaseAllowance definition follows the standard Defining decreaseAllowance input and output variables as standard	✓						
56	BP	The increaseAllowance definition follows the standard Defining increaseAllowance input and output variables as standard	✓						
57	BP	Minimize attack surface Checking whether all the external functions are necessary or not	✓	✓	✓				
58	BP	Transfer to the burn address is reverted Reverting transfer to 0x0 due to risk of total supply reduction	✓						
59	BP	Source code is decentralized Not using hard-coded addresses in the code	✓	✓					
60	BP	Funds can be held only by user-controlled wallets Transferring tokens to users to avoid creating a secondary market	!						
61	BP	Code logic is simple to understand Avoiding code nesting which makes the code less intuitive	✓	✓					
62	BP	All functions are documented Using NatSpec format to explain expected behavior of functions	✓						
63	BP	The Approval event is correctly logged Emitting Approval event in the approve() method	✓						
64	BP	Acceptable gas cost of the approve() function Checking for maximum 50000 gas cost when executing the approve()	!						
65	BP	Acceptable gas cost of the transfer() function Checking for maximum 60000 gas cost when executing the transfer()	!						
66	BP	Emitting event when state changes Emitting Change event when changing state variable values	✓						
67	BP	Use of unindexed arguments Using indexed arguments to facilitate external tools log searching		✓			✓	✓	
68	BP	ERC-20 compliance Implementing all 6 functions and 2 events as specified in EIP-20	✓	✓	✓		✓	✓	
69	BP	Conformance to naming conventions Following the Solidity naming convention to avoid confusion						✓	
70	BP	Token decimal Declaring token decimal for external apps when displaying balances	✓						
71	BP	Locked money (Freezing ETH) Implementing withdraw/reject functions to avoid ETH lost		✓			✓	✓	
72	BP	Malicious libraries Not using modifiable third-party libraries		✓					
73	BP	Payable fallback function Adding either fallback() or receive() function to receive ETH		✓			✓		
74	BP	Prefer external to public visibility level Improving the performance by replacing public with external		✓				✓	
75	BP	Token name Adding a token name variable for external apps	✓						
76	BP	Error information in revert condition Adding error description in require()/revert() to clarify the reason					✓		
77	BP	Complex Fallback Logging operations in the fallback() to avoid complex operations					✓		
78	BP	Function Order Following fallback, external, public, internal and private order					✓		
79	BP	Visibility Modifier Order Specifying visibility first and before modifiers in functions						✓	
80	BP	Non-initialized return value Not specifying return for functions without output		✓			✓		
81	BP	Token symbol Adding token symbol variable for usage of external apps	✓						
82	BP	Allowance spending is possible Ability of token transfer by transferFrom() to transfer tokens on behalf of another usercalc	✓						
99.5% success rate in performed audits by considering 'False Positives' and 'Informational' checks as 'Passed' (More details in section3)			100%	100%	100%	100%	100%	100%	97%

Table 3. Continuation of Table2.

clear from the tools’ own descriptions). Space will not permit us to discuss each one at the same level of detail as the ones we highlight in sections A and B, however we will include a simple statement describing the issue and the mitigation.

Since each tool employs different methodology to analyze smart contracts (e.g., comparing with violation patterns, applying a set of rules, using static analysis, *etc.*), there are false positives to manually check. Some false positives (e.g., *noReentrancy* modifier in *MythX* and *ContractGuard*) are not due to old/unmaintained rules. Other tools cannot reason adequately when **Modifiers** are used. The following are some examples of false positives (which we do not count in calculating our success rate):

1. *MythX* detects *Re-entrancy attack* in the *noReentrancy* modifier. In Solidity, modifiers are not like functions. They are used to add features or apply some restriction on functions[60]. Using modifiers is a known technique to implement Mutex and mitigate the attack[69]. This is a false positive and note that other tools have not identified the attack in modifiers.
2. *ContractGuard* flags *Re-entrancy attack* in **transfer()** function while both CEI and Mutex are implemented.
3. *Slither* detects two *low level call* vulnerabilities[35]. This is due to use of **call.value()** that is recommend way of transferring ETH after *Istanbul* hard-fork (EIP-1884). Therefore, adapting analyzers to new standards can improve accuracy of the security checks.
4. *SmartCheck* recommends not using **SafeMath** and check explicitly where overflows might be occurred. We consider this failed audit as false possible whereas utilizing **SafeMath** is a known technique to mitigate over/under flows. It also flags *using a private modifier* as a vulnerability by mentioning, “miners have access to all contracts’ data and developers must account for the lack of privacy in Ethereum”. However private visibility in Solidity concerns object oriented inheritance not confidentiality. For actual confidentiality, the best practice is to encrypt private data or store them off-chain. The tool also warns against **approve()** in ERC-20 due to *front-running attacks*. Despite EIP-1884, it still recommends using of **transfer()** method with stipend of 2300 gas. There are other false positives such as SWC-105 and SWC-112 that are passed by other tools.
5. *Securify* detects the *Re-entrancy* attack due to unrestricted writes in the **noReentrancy** modifier[68]. Modifiers are the recommended approach and are not accessible by users. It also flags *Delegatecall to Untrusted Callee* (SWC-112) while there is no usage of **delegatecall()** in the code. It might be due to use of **SafeMath** library which is an embedded library. In Solidity, embedded libraries are called by JUMP commands instead of **delegatecall()**. Therefore, excluding embedded libraries from this check might improve accuracy of the tool. Similar to *SmartCheck*, it still recommends to use the **transfer()** method instead of **call.value()**.
6. *EY token review* considers **decreaseAllowance** and **increaseAllowance** as standard ERC-20 functions and if not implemented, recognizes the code as vulnerable to a *front-running*. These two functions are not defined in the

ERC-20 standard[23] and considered only by this tool as mandatory functions. There are other methods to prevent the attack while adhering ERC-20 specifications (see Rahimian *et al.* for a full paper on this attack and the basis of the mitigation in **TokenHook** [51]). The tool also falsely detects the *Overflow*, mitigated through **SafeMath**. Another identified issue is *Funds can be held only by user-controlled wallets*. The tool warns against any token transfer to Ethereum addresses that belong to smart contracts. However, interacting with ERC-20 token by other smart contracts was one of the main motivations of the standard. It also checks for maximum 50000 gas in `approve()` and 60000 in `transfer()` method. We could not find corresponding SWC registry or standard recommendation on these limitations and therefore consider them as informational.

7. *Odin* raises *Outdated compiler version* issue due to locking solidity version to 0.5.11. We have used this version due to its compatibility with other auditing tools. Furthermore, other tools have not identified such an issue and we therefore consider it as informational.

3.2 Comparing audits

After manually overriding the false positives, the average percentage of passed checks for **TokenHook** reaches to 99.5%. To pass the one missing check and reach a 100% success rate across all tools, we prepared the same code in Solidity version 0.8.0, however it cannot be audited anymore with most of the tools.

We repeated the same auditing process on the top ten tokens based on their market cap [21]. The result of all these evaluation have been summarized in Table 4 by considering false positives as failed audits. This provide the same evaluation conditions across all tokens. Since each tool uses different analysis methods, number of occurrences are considered for comparisons. For example, **MythX** detects two *re-entrancy* in **TokenHook**; therefore, two occurrences are counted instead of one.

As it can be seen in Table 4, **TokenHook** has the least number of security flaws (occurrences) compared to other tokens. We stress that detected security issues for **TokenHook** are all false positives. We are also up-front that this metric is not a perfect indication of security. The other tokens may also have many/all false positives (such an analysis would be interesting future work), and not all true positives can be exploited [47]. Mainly, we want to show this measurement as being consistent with our claims around the security of **TokenHook**. Had **TokenHook**, for example, had the highest number of occurrences, it would be a major red flag.

4 Conclusion

98% of tokens on Ethereum today implement ERC-20. While attention has been paid to the security of Ethereum DApps, threats to tokens can be specific to ERC-20 functionality. Further, there is no vulnerability reference site (*cf.* the

ERC-20 Token	Auditing Tool							Total issues
	EY Token Review	Smart Check	Securify	MythX (Mythril)	Contract Guard	Slither	Odin	
TokenHook	9	11	4	2	10	2	2	40
TUSD	20	11	2	1	14	16	6	70
PAX	16	9	6	4	16	13	9	73
USDC	17	9	6	5	18	15	10	80
INO	11	10	14	8	14	24	12	93
HEDG	10	28	11	1	29	24	16	119
BNB	13	21	12	13	41	39	3	142
MKR	11	27	38	9	16	34	18	153
LINK	12	27	38	9	16	34	18	181
USDT	12	29	8	17	46	55	30	197
LEO	32	25	8	23	70	75	19	252

Table 4. Security flaws detected by seven auditing tools in **TokenHook** (the proposal) compared to top 10 ERC-20 tokens by market capitalization in May 2020. **TokenHook** has the lowest reported security issues (occurrences).

SWC Registry) specifically for ERC-20 tokens. In this paper, we provide a detailed study of ERC-20 security, collecting and deduplicating 82 vulnerabilities and best practices, examining the ability of seven audit tools, and auditing 10 ERC-20 deployments. Most importantly, we provide a concrete implementation of ERC-20 called **TokenHook**. It is designed to be secure against known vulnerabilities, and can serve as a second reference implementation to provide software diversity. We test it at Solidity version 0.5.11 (due to the limitation of the audit tools) and also provide it at 0.8.0. **TokenHook** can be used as template to deploy new ERC-20 tokens, migrate current vulnerable deployments, and to benchmark the precision of Ethereum audit tools.

References

1. Alex Beregszaszi, A.S.: Hardfork meta: Istanbul. <https://github.com/ethereum/EIPs/blob/master/EIPS/eip-1679.md> (Dec 2019)
2. di Angelo, M., Salzer, G.: A survey of tools for analyzing ethereum smart contracts. https://publik.tuwien.ac.at/files/publik_278277.pdf (Aug 2019)
3. Breidenbach, L., Daian, P., Juels, A., Gun Sirer, E.: An in-depth look at the parity multisig bug. <https://hackingdistributed.com/2017/07/22/deep-dive-parity-bug/> (Jul 2017)
4. Bulgakov, K.: Three methods to send ether by means of solidity. <https://medium.com/daox/three-methods-to-transfer-funds-in-ethereum-by-means-of-solidity-5719944ed6e9> (Feb 2018)
5. Chen, T., Zhang, Z., Li, Z.: Tokenscope: Automatically detecting inconsistent behaviors of cryptocurrency tokens in ethereum. <http://www4.comp.polyu.edu.hk/~csxluo/TokenScope.pdf> (Nov 2019)
6. Clark, J., Bonneau, J., Miller, A., Kroll, J.A., Felten, E.W., Narayanan, A.: On decentralizing prediction markets and order books. In: WEIS (2014)
7. ConsenSys: Mythx swc coverage. <https://mythx.io/swc-coverage/> (Nov 2019)
8. ConsenSys: Tokens. <https://github.com/ConsenSys/Tokens> (Jun 2020)

9. Coverdale, C.: Transaction-ordering attacks. <https://medium.com/coinmonks/solidity-transaction-ordering-attacks-1193a014884e> (Mar 2018)
10. Cremer, L.: Missing return value bug — at least 130 tokens affected. <https://medium.com/coinmonks/missing-return-value-bug-at-least-130-tokens-affected-d67bf08521ca> (Jun 2018)
11. Deployer, B.: Beautychain (bec). <https://etherscan.io/address/0xc5d105e63711398af9bbff092d4b6769c82f793d> (Jun 2020)
12. Diligence, C.: Ethereum smart contract security best practices. <https://consensys.github.io/smart-contract-best-practices/> (Jan 2020)
13. Diligence, C.: Token implementation best practice. <https://consensys.github.io/smart-contract-best-practices/tokens/> (Mar 2020)
14. documentation, E.: Re-entrancy. <https://solidity.readthedocs.io/en/latest/security-considerations.html#re-entrancy> (Jan 2020)
15. documentation, S.: Security considerations. <https://solidity.readthedocs.io/en/latest/security-considerations.html> (Jan 2020)
16. Durov, P.: What was ton and why it is over. <https://telegra.ph/What-Was-TON-And-Why-It-Is-Over-05-12> (May 2020)
17. Eskandari, S., Moosavi, S., Clark, J.: Sok: Transparent dishonesty: front-running attacks on blockchain. *International Conference on Financial Cryptography and Data Security* **1**, 380 (2019)
18. Ethereum: Project repository. <https://github.com/ethereum> (May 2014)
19. Ethereum: Solidity — solidity documentation. <https://solidity.readthedocs.io/en/latest/contracts.html?highlight=library#libraries> (Jan 2020)
20. Ethereum: Solidity — solidity documentation. <https://solidity.readthedocs.io/en/latest/> (Jan 2020)
21. EtherScan: Token tracker. <https://etherscan.io/tokens?sortcmd=remove&sort=marketcap&order=desc> (Apr 2020)
22. EY: Token review. <https://review-tool.blockchain.ey.com> (Sep 2019)
23. Fabian Vogelsteller, V.B.: Erc-20 token standard. <https://github.com/ethereum/EIPs/blob/master/EIPS/eip-20.md> (Nov 2015)
24. Ferreira Torres, C., Schutte, J., State, R.: Osiris: Hunting for integer bugs in ethereum smart contracts. <https://dl.acm.org/doi/10.1145/3274694.3274737> (Dec 2018)
25. Finley, K.: A \$50 million hack just showed that the dao was all too human — wired. <https://www.wired.com/2016/06/50-million-hack-just-showed-dao-human/> (Sep 2016)
26. Forrest, S., Somayaji, A., Ackley, D.H.: Building diverse computer systems. In: *Proceedings. The Sixth Workshop on Hot Topics in Operating Systems (Cat. No.97TB100133)*. pp. 67–72 (1997). <https://doi.org/10.1109/HOTOS.1997.595185>
27. Forrest, S., Hofmeyr, S.A., Somayaji, A.: Computer immunology. *Commun. ACM* **40**(10), 88–96 (Oct 1997). <https://doi.org/10.1145/262793.262811>, <https://doi.org/10.1145/262793.262811>
28. GmbH, G.F.: Golem network token. <https://etherscan.io/address/0xa74476443119A942dE498590Fe1f2454d7D4aC0d#code> (Nov 2016)
29. GuardStrike: Contractguard knowledge-base. <https://contract.guardstrike.com/#/knowledge> (Mar 2020)
30. Hale, T.: Resolution on the eip20 api approve / transferfrom multiple withdrawal attack #738. <https://github.com/ethereum/EIPs/issues/738> (Oct 2017)
31. Holst Swende, M.: Repricing for trie-size-dependent opcodes. <https://github.com/ethereum/EIPs/blob/master/EIPS/eip-1884.md> (Mar 2019)

32. Huashan Chen, Marcus Pendleton, L.N., Xu, S.: A survey on ethereum systems security: Vulnerabilities, attacks and defenses. <https://arxiv.org/pdf/1908.04507.pdf> (Aug 2019)
33. Jain, S.: All you should know about libraries in solidity. <https://medium.com/coinmonks/all-you-should-know-about-libraries-in-solidity-dd8bc953eae7> (Sep 2018)
34. Ji, R., He, N., Wu, L., Wang, H.: Deposafe: Demystifying the fake deposit vulnerability. <https://arxiv.org/pdf/2006.06419.pdf> (Jun 2020)
35. Josselin, F.: Slither – a solidity static analysis framework. <https://blog.trailofbits.com/2018/10/19/slither-a-solidity-static-analysis-framework/> (Oct 2018)
36. Josselin, F.: Slither – detector documentation. <https://github.com/crytic/slither/wiki/Detector-Documentation#name-reused> (Mar 2020)
37. Lando, G.: Guy lando’s knowledge list. <https://github.com/guylando/KnowledgeLists/blob/master/EthereumSmartContracts.md> (May 2019)
38. Luu, L., Chu, D.H., Olickel, H., Saxena, P., Hobor, A.: Making smart contracts smarter. https://dl.acm.org/ft_gateway.cfm?id=2978309&ftid=1805715&dwn=1&CFID=86372769&CFTOKEN=b697c89273876526-8CBDF39B-A89A-31D2-F565B24919F796C6 (Oct 2016)
39. Manning, A.: Comprehensive list of known attack vectors and common anti-patterns. <https://github.com/sigp/solidity-security-blog> (Nov 2019)
40. Marx, S.: Stop using solidity’s transfer() now. <https://diligence.consensys.net/blog/2019/09/stop-using-soliditys-transfer-now/> (Dec 2019)
41. OpenZeppelin: Proxy patterns. <https://blog.openzeppelin.com/proxy-patterns/> (Apr 2018)
42. OpenZeppelin: Contracts. <https://github.com/OpenZeppelin/openzeppelin-contracts/blob/master/contracts/token/ERC20/ERC20.sol> (Jun 2020)
43. OpenZeppelin: Contracts. <https://github.com/OpenZeppelin/openzeppelin-contracts/blob/master/contracts/math/SafeMath.sol> (Jun 2020)
44. Palladino, S.: The parity wallet hack explained. <https://blog.openzeppelin.com/on-the-parity-wallet-multisig-hack-405a8c12e8f7/> (July 2017)
45. Parity: Security alert. <https://www.parity.io/security-alert-2/> (Nov 2018)
46. PeckShield: Alert: New batchoverflow bug in multiple erc20 smart contracts. <https://blog.peckshield.com/2018/04/22/batchOverflow/> (Apr 2018)
47. Perez, D., Livshits, B.: Smart contract vulnerabilities: Vulnerable does not imply exploited. https://www.usenix.org/system/files/sec21summer_perez.pdf (Aug 2018)
48. Quintal, J.: Building robust smart contracts with openzeppelin. <https://www.trufflesuite.com/tutorials/robust-smart-contracts-with-openzeppelin> (Aug 2017)
49. Qureshi, H.: A hacker stole \$31m of ether — how it happened, and what it means. <https://www.freecodecamp.org/news/a-hacker-stole-31m-of-ether-how-it-happened-and-what-it-means-for-ethereum-9e5dc29e33ce/> (Jul 2017)
50. Rahimian, R.: Overflow attack in ethereum smart contracts. <https://blockchain-projects.readthedocs.io/overflow.html> (Dec 2018)
51. Rahimian, R., Eskandari, S., Clark, J.: Resolving the multiple withdrawal attack on erc20 tokens. <https://arxiv.org/abs/1907.00903> (Jul 2019)
52. Reinno: A new way to own and invest in real estate. <https://reinno.io/tokenization.html> (Jun 2020)
53. Reporting, A.: Generate meaningful knowledge from ethereum. <https://reports.aeth.io/> (Jul 2020)

54. Ritzdorf, H.: Ethereum istanbul hardfork, the security perspective. <https://docs.google.com/presentation/d/1IiRYSjwle02zQUmWId06Bss8GrxGyw6nQAiZdCRFEpk/> (Oct 2019)
55. Roan, A.: 7 openzeppelin contracts you should always use. <https://medium.com/better-programming/7-openzeppelin-contracts-you-should-always-use-5ba2e7953cc4> (Apr 2020)
56. Rodler, M., Li, W., Karame, G., Davi, L.: Sereum: Protecting existing smart contracts against re-entrancy. <https://arxiv.org/pdf/1812.05934.pdf> (Dec 2018)
57. Rubixi: Rubixi contract. <https://etherscan.io/address/0xe82719202e5965Cf5D9B6673B7503a3b92DE20be#code> (Mar 2016)
58. Sedgwick, K.: Myetherwallet servers are hijacked. <https://news.bitcoin.com/myetherwallet-servers-are-hijacked-in-dns-attack/> (Apr 2018)
59. Shirshova, N.: The benefits of “buy” and “sell”. <https://medium.com/orbise/the-benefits-of-buy-and-sell-token-functions-dcea536aaf7c> (Dec 2018)
60. Simon, G.: Solidity modifier tutorial - control functions with modifiers. <https://coursetro.com/posts/code/101/Solidity-Modifier-Tutorial---Control-Functions-with-Modifiers> (Oct 2017)
61. SmartContractSecurity: Smart contract weakness classification and test cases. <https://swcregistry.io/> (Jun 2020)
62. SmartDec: Knowledge-base. <https://tool.smartdec.net/knowledge> (Sep 2018)
63. Sooho: Verify a smart contract. <https://odin.sooheo.io/> (Mar 2020)
64. Spagnuolo, F.: Ethereum in depth, part 2. <https://blog.openzeppelin.com/ethereum-in-depth-part-2-6339cf6bddb9/> (Jul 2018)
65. Tanner, J.: Summary of ethereum upgradeable smart contract r&d — part 1–2018. <https://blog.indorse.io/ethereum-upgradeable-smart-contract-strategies-456350d0557c> (Mar 2018)
66. Tracker, T.: Erc-20 tokens. <https://etherscan.io/tokens> (Jul 2020)
67. Tsankov, P.: Securify v2.0. <https://github.com/eth-sri/securify2> (Jan 2020)
68. Tsankov, P., Dan, A., Drachsler-Cohen, D., Gervais, A.: Securify: Practical security analysis of smart contracts. <https://arxiv.org/pdf/1806.01143.pdf> (Aug 2018)
69. Venturo, N., Giordano, F.: Reentrancy guard. <https://github.com/OpenZeppelin/openzeppelin-contracts/blob/master/contracts/utils/ReentrancyGuard.sol> (Oct 2017)
70. Vessenes, P.: Monolithdao. <https://github.com/MonolithDAO/token/blob/master/src/Token.sol> (Apr 2017)
71. Vladimirov, M.: Attack vector on erc20 api (approve/transferfrom methods) and suggested improvements. <https://github.com/ethereum/EIPs/issues/20#issuecomment-263524729> (Nov 2016)
72. wikipedia: Mutex. en.wikipedia.org/wiki/Mutual_exclusion (Jan 2019)
73. Wikipedia: Liberty reserve. https://en.wikipedia.org/wiki/Liberty_Reserve (Jun 2020)
74. Wood, G.: Ethereum: a secure decentralised generalised transaction ledger. <http://gavwood.com/paper.pdf> (Mar 2016)
75. Zalecki, M.: Create and distribute your erc20 token with openzeppelin. <https://www.tooploox.com/blog/create-and-distribute-your-erc20-token-with-openzeppelin> (Aug 2018)

A Sample of high profile vulnerabilities

In this section, we sample some high profile vulnerabilities, typically ones that have been exploited in real world ERC-20 tokens[39,32,15,12,37]. For each, we

- (i) briefly explain technical details, (ii) the ability to affect ERC-20 tokens, and
- (iii) discuss mitigation techniques.

A.1 Multiple withdrawal

This ERC-20-specific issue was originally raised in 2017[71,30]. It can be considered as a *transaction-ordering*[9] or *front-running*[17] attack. There are two ERC-20 functions (*i.e.*, `Approve()` and `transferFrom()`) that can be used to authorize a third party for transferring tokens on behalf of someone else. Using these functions in an undesirable situation (*i.e.*, front-running or race-condition) can result in allowing a malicious authorized entity to transfer more tokens than the owner wanted. There are several suggestions to extend ERC-20 standard (*e.g.*, MonolithDAO[70] and its extension in OpenZeppelin[42]) by adding new functions (*i.e.*, `decreaseApproval()` and `increaseApproval()`), however, securing `transferFrom()` method is the effective one while adhering specifications of the ERC-20 standard[51].

A.2 Arithmetic Over/Under Flows.

An *integer overflow* is a well known issue in many programming languages. For ERC-20, one notable exploit was in April 2018 that targeted the BEC Token[11] and resulted in some exchanges (*e.g.*, OKEx, Poloniex, *etc.*) suspending deposits and withdrawals of all tokens. Although BEC developers had considered most of the security measurements, only line 261 was vulnerable[24][46]. The attacker was able to pass a combination of input values to transfer large amount of tokens[50]. It was even larger than the initial supply of the token, allowing the attacker to take control of token financing and manipulate the price. In Ethereum, integer overflows do not throw an exception at runtime. This is by design and can be prevented by using the `SafeMath` library[43] wherein `a+b` will be replaced by `a.add(b)` and throws an exception in the case of arithmetic overflow.

A.3 Re-entrancy

One of the most studied vulnerabilities is re-entrancy, which resulted in a US\$50M attack on a DApp (called the DAO) in 2016 and triggered an Ethereum hard-fork to revert[25]. At first glance, re-entrancy might seem inapplicable to ERC-20 however any function that changes internal state, such as balances, need to be checked. Further, some ERC-20 extensions could also be problematic. One example is ORBT tokens [52] which support token exchange with ETH without going through a crypto-exchange [59]: an attacker can call the exchange function to sell the token and get back equivalent in ETH. However, if the ETH is transferred in a vulnerable way before reaching the end of the function and updating the balances, control is transferred to the attacker receiving the funds and the same function could be invoked over and over again within the limits of a single transaction, draining excessive ETH from the token contract. This

variant of the attack is known as *same-function re-entrancy*, but it has three other variants: *cross-function*, *delegated* and *create-based* [56]. Mutex [72] and CEI [14] techniques can be used to prevent it. In Mutex, a state variable is used to lock/unlock transferred ETH by the lock owner (*i.e.*, token contract). The lock variable fails subsequent calls until finishing the first call and changing requester balance. CEI updates the requester balance before transferring any fund. All interactions (*i.e.*, external calls) happen at the end of the function and prevents recursive calls. Although CEI does not require a state variable and consumes less Gas, developers must be careful enough to update balances before external calls. Mutex is more efficient and blocks *cross-function* attacks at the beginning of the function regardless of internal update sequences. CEI can also be considered as a best practice and basic mitigation for the *same-function re-entrancy*

A.4 Unchecked return values

In Solidity, sending ETH to external addresses is supported by three options: `call.value()`, `transfer()`, or `send()`. The `transfer()` method reverts all changes if the external call fails, while the other two return a boolean value and manual check is required to revert transaction to the initial state[4]. Before the *Istanbul* hard-fork[1], `transfer()` was the preferred way of sending ETH. It mitigates reentry by ensuring ETH recipients would not have enough gas (*i.e.*, a 2300 limit) to do anything meaningful beyond logging the transfer when execution control was passed to them. EIP-1884[31] has increased the gas cost of some opcodes that causes issues with `transfer()`⁴. This has led to community advice to use `call.value()` and rely on one of the above re-entrancy mitigations (*i.e.*, Mutex or CEI)[40,54].

A.5 Frozen Ether

As ERC-20 tokens can receive and hold ETH, just like a user accounts, functions need to be defined to withdraw deposited ETH (including unexpected ETH). If these functions are not defined correctly, an ERC-20 token might hold ETH with no way of recovering it (*cf.* Parity Wallet[45]). If necessary, developers can require multiple signatures to withdraw ETH.

A.6 Unprotected Ether Withdrawal

Improper access control may allow unauthorized persons to withdraw ETH from smart contracts (*cf.* Rubixi[57]). Therefore, withdrawals must be triggered by only authorized accounts and ideally multiple parties.

⁴ After *Istanbul*, the `fallback()` function consumes more than 2300 Gas if called via `transfer()` or `send()` methods.

A.7 State variable manipulation

The `DELEGATECALL` opcode enables a DApp to invoke external functions of other DApps and execute them in the context of calling contract (*i.e.*, the invoked function can modify the state variables of the caller). This makes it possible to deploy libraries once and reuse the code in different contracts. However, the ability to manipulate internal state variables by external functions has led to incidents where the entire contract was hijacked (*cf.* the second hack of Parity MultiSig Wallet[3]). Preventive techniques is to use `Library` keyword in Solidity to force the code to be stateless, where data is passed as inputs to functions and passed back as outputs and no internal storage is permitted[19]. There are two types of Library: *Embedded* and *Linked*. Embedded libraries have only internal functions (EVM uses `JUMP` opcode instead of `DELEGATECALL`), in contrast to linked libraries that have public or external functions (EVM initiate a “message call”). Deployment of linked libraries generates a unique address on the blockchain while the code of embedded libraries will be added to the contract’s code [33]. It is recommended to use Embedded libraries to mitigate this attack.

A.8 Public visibility

In Solidity, visibility of functions are `Public` by default and they can be called by any external user/contract. In the Parity MultiSig Wallet hack[49], an attacker was able to call public functions and reset the ownership address of the contract, triggering a \$31M USD theft. It is recommended to explicitly specify visibility of functions instead of default `Public` visibility.

B A sample of best practices

Due to space, we highlight a few best practices that have been accepted by the Ethereum community to proactively prevent known vulnerabilities[13]. Some best practices are specific to ERC-20, while others are generic for all DApps—in which case, we discuss their relevance to ERC-20.

B.1 Compliance with ERC-20.

According to the ERC-20 specifications, all six methods and two events must be implemented and are not optional. Tokens that do not implement all methods (*e.g.*, GNT which does not implement the `approve()`, `allowance()` and `transferFrom()` functions due to *front-running*[28]) can cause failed function calls from other applications. They might also be vulnerable to complex attacks (*e.g.*, Fake deposit vulnerability[34], Missing return value bug[10]).

B.2 External visibility.

Solidity supports two types of *function calls*: internal and external [20]. Note that functions calls are different than functions visibility (*i.e.*, `Public`, `Private`,

Internal and External) which confusingly uses overlapping terminology. Internal function calls expect arguments to be in memory and the EVM copies the arguments to memory. Internal calls use `JUMP` opcodes instead of creating an *EVM call*.⁵ Conversely, External function calls create an *EVM call* and can read arguments directly from the `calldata` space. This is cheaper than allocating new memory and designed as a read-only byte-addressable space where the data parameter of a transaction or call is held[64]. A best practice is to use external visibility when we expect that functions will be called externally.

B.3 Fail-Safe Mode.

In the case of a detected anomaly or attack on a deployed ERC-20 token, the functionality of the token can be frozen pending further investigation. For regulated tokens, the ability for a regulator to issue a ‘cease trade’ order is also generally required.

B.4 Firing events.

In ERC-20 standard, there are two defined events: **Approval** and **Transfer**. The first event logs successful allowance changes by token holders and the second logs successful token transfers by the `transfer()` and `transferFrom()` methods. These two events must be fired to notify external application on occurred changes. The external application (*e.g.*, TokenScope[5]) might use them to detect inconsistent behaviors, update balances, show UI notifications, or to check new token approvals. It is a best practice to fire an event for every state variable change.

B.5 Global or Miner controlled variables.

Since malicious miners have the ability to manipulate global Solidity variables (*e.g.*, `block.timestamp`, `block.number`, `block.difficulty`, *etc.*), it is recommended to avoid these variables in ERC-20 tokens.

B.6 DoS with Unexpected revert.

A function that attempts to complete many operations that individually may revert could deadlock if one operation always fails. For example, `transfer()` can throw an exception—if one transfer in a sequence fails, the whole sequence fails. One standard practice is to account for ETH owed and require withdrawals through a dedicated function. In `TokenHook`, ETH is only transferred to a single party in a single function `sell()`. It seems overkill to implement a whole accounting system for this. As a consequence, a seller that is incapable of receiving ETH (*e.g.*, operating from a contract that is not payable) will be unable to sell their tokens for ETH. However they can recover by transferring the tokens to a new address to sell from.

⁵ Also known as “message call” when a contract calls a function of another contract.